

# CMPT 756.203 T4-Wednesday Milestone 2

Team	T4-Wednesday		
Section	G100		
Team Member Name	Kunal Niranjana Desai	Rohan Harode	Juan Ospina
Email (SFU)	kdesai@sfu.ca	rharode@sfu.ca	jospina@sfu.ca
Email (G-Suite/Gmail)	kdesai1505@gmail.com	harode48@gmail.com	juandavidospina@gmail.com
Github ID	kunaldesai97	rohanharode	juandavidospina
Additional Notes			

# Section 1 - Problem Statement

## Problem Domain & Application

### **Online Bill Payment System: "My Bill Payments"**

Our client, a recognized Fintech named "ABCFintech," wants to offer its current and potential customers an application to pay their bills online. This application is called "My Bill Payments". Through this application, a customer gets a centralized platform to pay all their bills from different companies like BC Hydro (Electricity), Rogers (cellphone), Telus (Internet) (from now on, named "billers") as well as manage them, thus saving time and avoiding overdue balances and late fees.

For example, a person named John Doe wants to pay three bills at the end of the month: BC Hydro, Cellphone (Rogers), and Internet (Telus). John will be able to log in on our application, and through the "Bill retrieval" functionality, he can look up the details of each bill, using as input the biller name and the bill identifier, e.g., Biller name: "Rogers" and Bill ID: 12345. Assuming that the bill exists, our system will return the bill's relevant information: Biller name, Bill identifier, amount to pay, and due date. Then, John can pay this bill using his credit card. Hence, John will enter its credit card information into the system: credit card number (16 digits), cardholder name, expiration year-month, and the three-digit security code. In case that John has funds available to cover this payment, our system will flag the bill as paid (preventing that it can be paid again by another user) and register the payment information (credit card first four and last four digits, cardholder name and expiration year-month). The user can repeat this process as many times as the number of bills. Also, if John wants to review his paid bills, our system will provide a functionality to review all his past payments (payments history).

# Specification of REST API (Microservices Contract)

Version: v1  
Service: Gateway  
Visibility: Public  
Domain: Ingress-gateway  
Serialized Data/Content-Type: json

API	Description
/api/v1/user	Users service URL
/api/v1/bill	Bills service URL
/api/v1/billers	Billers service URL

## Service: Users

Version: v1  
Visibility: Private  
Domain: Users  
Serialized Data/Content-Type: json

API	Description	Request Body/Parameters	Response Body	HTTP Response Code	Error Codes	Request Example	Response Example

PUT - /api/v1/user/login	Login	JSON:{"uname": user_name, "pword": password}	Hash of user context suitable for passing to other calls	200	500	PUT https://host:5000/api/v1/user/login	{"UserContext": ""}
PUT - /api/v1/user/logout	Logout	JSON: {"jwt": token}	None	200	500	PUT https://host:5000/api/v1/user/logout	{ Message: ok }
PUT - /api/v1/user/	UPDATE one user	Body: { uname: string, pword: string, Fname: string, Lname: string, Email: string, Secquestion: string, Secanswer:string }	OK response	200	500	PUT https://host:5000/api/v1/user/	{ Message: ok }
POST - /api/v1/user/	CREATE one user	Body: { uname: string, Fname: string, Lname: string, Email: string, Secquestion:st	OK response	200	500	POST https://host:5000/api/v1/user	{ ResponseMetadata: { ... etc. } }

		ring, Secanswer:string }  Params: None					
DELETE - /api/v1/user/	DELETE one user by username		JSON of response from aws	200	500	DELETE https://host:5000/api/v1/user/	{ ResponseMetadata: { ... etc. } }
GET - /api/v1/user/< uname>	GET one user by username		JSON of User entity	200	500	GET https://host:5000/api/v1/user?user=johndoe	{User Object}
GET - /api/v1/user?u serkey	GET user by key	Query Params: userkey:string	JSON of User entity	200	500	GET https://host:5000/api/v1/user?userkey=email:johndoe@gmail.com	{User Object}

## Service: Bill

Version: v1

Visibility: Private

Domain: Bills

Serialized Data/Content-Type: json

API	Description	Request Body/Parameters	Response Body	HTTP Response Code	Error Codes	Request Example	Response Example
GET - /api/v1/bill/	Retrieve one bill	Param: bill_id	None	200	500	GET https://host:5001/api/v1/bill/	{ bill_id: string, biller_id:string, amount:float, due_date: date, bill_paid:boolean}
POST - /api/v1/bill/	Insert one bill	Body: { user_id: string, biller_id:string, bill_amount:float, due_date:date, bill_paid: boolean (default: False)} User-context token	OK response	200	500	POST https://host:5001/api/v1/bill/	{ Message: ok }

DELETE - /api/v1/bill/	DELETE one bill	Param: bill_id	JSON of response from aws	200	500	DELETE https://host:5001/api/v1/bi ll/	{ ResponseMetada ta: { ... etc. } }
PUT - /api/v1/bill/ pay/<bill_id >	Pay bill	Body: {“cc_number”:, cc_exp_dat: String, payment_date: String}  Header: User-context token	OK response	200	500	PUT https://host:5000/api/v1/bi ll/pay	{ Message: ok }
GET - /api/v1/bill/ <uname>	GET bills by username	Query Params:  username:string	JSON of Bill entity	200	500	GET https://host:5000/api/v1/bi ll/johndoe	{Bill Object}

## Service: Biller

Version: v1

Visibility: Private

Domain: Billers

Serialized Data/Content-Type: json

API	Description	Request Body/Parameters	Response Body	HTTP Response Code	Error Codes	Request Example	Response Example
GET - /api/v1/biller/	Retrieve one biller	Param: biller_id	None	200	500	GET https://host:5001/api/v1/biller/	{ biller_id: string, biller_name:string, biller_description:string}
POST - /api/v1/biller	Insert one biller	Body: { biller_id: string, biller_name: string, biller_description:string }	OK response	200	500	POST https://host:5001/api/v1/biller/	{ Message: ok }
DELETE - /api/v1/biller/	DELETE one biller	Param: biller_id	JSON of response from aws	200	500	DELETE https://host:5001/api/v1/biller/	{ ResponseMetadata: { ... etc. } }



## Service: db

Version: v1

Visibility: Private

Domain: Datastore

Serialized Data/Content-Type: json

API	Description	Request Body/Parameters	Response Body	HTTP Response Code	Error Code	Request Example	Response Example
GET - /api/v1/datastore/read	Read an object	Param: object-type, object-key	JSON of aws response	200	500	GET https://host:5002/api/v1/datastore/read?objtype=user&objkey=johndoe	{ "Count": 1, "Items": [ { "uname": "johndoe", "fname": "John", "lname": "Doe", "email": "johndoe@gmail.com", "secquestion": "Name of your first pet?", "secanswer": "Lucky" } ], "ResponseMetadata": { "HTTPHeaders": { "connection": "keep-alive", "content-length": "165", "content-type": "application/x-amz-json-1.0", "date": "Sat, 12 Sep 2020 18:16:15 GMT", "server": "Server", "x-amz-crc32": "196980578", "x-amzn-requestid": "AOGKN903DF66VLU3GEBEE08DK3VV4KQNS05AEMVJF66Q9ASUAAJG" }, "HTTPStatusCode": 200, "RequestId": "AOGKN903DF66VLU3GEBEE08DK3VV4K"

							QNSO5AEMVJF66Q9ASUAAJG", "RetryAttempts": 0 }, "ScannedCount": 1 }
POST - /api/v1/datastore/write	Write an object	Body: objtype, object-key(s)	ID of new entity	200	500	POST https://host:5002/api/v1/datastore/write	{ "uname": "johndoe" }
DELETE = /api/v1/datastore/delete	Delete an object	Param: objtype, object-key	JSON of aws response	200	500	DELETE https://host:5000/api/v1/datastore/delete?objtype=user&objkey=johndoe	{ "ResponseMetadata": { "HTTPHeaders": { "connection": "keep-alive", "content-length": "2", "content-type": "application/x-amz-json-1.0", "date": "Sat, 12 Sep 2020 18:13:04 GMT", "server": "Server", "x-amz-crc32": "2745614147", "x-amzn-requestid": "N7R6LO93FFBDH1A5GRRL55LBS7VV4KQNSO5AEMVJF66Q9ASUAAJG" }, "HTTPStatusCode": 200, "RequestId": "N7R6LO93FFBDH1A5GRRL55LBS7VV4KQNSO5AEMVJF66Q9ASUAAJG", "RetryAttempts": 0 } }
PUT - /api/v1/datastore/update	Update an object	Params: objtype, objkey	JSON of aws response	200	500	PUT https://host:5002/api/v1/datastore/update?objtype=user&objkey=email:johndoenew@gmail.com	{ "ResponseMetadata": { "HTTPHeaders": { "connection": "keep-alive", "content-length": "2", "content-type": "application/x-amz-json-1.0", "date": "Sat, 12 Sep 2020 18:13:04 GMT", "server": "Server", "x-amz-crc32": "2745614147", "x-amzn-requestid": "N7R6LO93FFBDH1A5GRRL55LBS7VV4K

		Body: object-key(s)  Example: email				QNSO5AEMVJF66Q9ASUAAJG" }, "HTTPStatusCode": 200, "RequestId": "N7R6LO93FFBDH1A5GRRL55LBS7VV4K QNSO5AEMVJF66Q9ASUAAJG", "RetryAttempts": 0 } }
--	--	---	--	--	--	--

## Database Schema (DynamoDB)

Table: Users

Tag	Value	Comment
lname	string	Last name of user
fname	string	First name of user
uname	string	Email ID of user
password	string	User's account password
security_question	string	Security question in case user forgets password

security_answer	string	Answer to the security question
user_id	string	Unique id generated by DynamoDB

Table: Bills

Tag	Value	Comment
bill_amount	decimal	Amount billed
bill_paid	boolean	true/false
due_date	string	Due date for the bill
payment_date	string	Date of the bill payment
cc_first_four_digits	integer	Identification of credit card franchise: Visa, MasterCard etc
cc_last_four_digits	integer	Credit card information for troubleshooting
cc_exp_date	string	Expiry date of credit card
user_id	string	User id
biller_id	string	Unique id generated by DynamoDB
bill_id	string	Unique id generated by DynamoDB

Table: Billers

Tag	Value	Comment
biller_name	string	Name of the billing company
active	boolean	true/false
biller_description	string	Description of the billing company
biller_id	string	Unique id not interpreted by DynamoDB

## Section 2 - Github Repo Guide

Path	Note
/docs	Information of various services
/docs/user_service	Details of users service and functionalities
/docs/bill_service	Details of bills service and functionalities
/docs/biller_service	Details of billers service and functionalities
/docs/minikube	Minikube steps, commands
/docs/docker	Docker steps, commands
/docs/AKS	Azure Kubernetes commands
/code	Code for API calls of different services
/code/users	Code for API calls of users services

/code/bills	Code for API calls of bills services
/code/billers	Code for API calls of billers services
/laC	Infrastructure as a Code directory
/laC/cloudformation	Cloudformation stacks
/laC/cluster	Cluster creation
/laC/k8s	Service gateway

## Section 3 - Reflection on Development

We adopted the idea of creating a product backlog to have a clear understanding of all the services we plan to implement for the project. We also had weekly meetings for 30 minutes each to catch up on the progress that will help us stay on track with our project plans, collaborate with others and brainstorm regarding the application and REST APIs. Since we are a team of 3 (with no scrum master or product owner), each of us took responsibility to make sure we were focusing on the bigger picture of the project while undertaking tasks. We have already started the Scrum process by using a Kanban dashboard on Github where we are creating stories and tasks. We are also maintaining a shared Google Doc to suggest and incorporate changes to the proposed application, architecture, REST API's and resolving them to promote a continuous cycle of improvement.

The readings which were most helpful to us were R2 (Scrum guide), R3 (Epics, stories, tasks, subtasks), and R6 (Revision control systems). The scrum guide gave us a good understanding of how scrum works and methodologies. It motivated us to try it for the term project starting with the product backlog and weekly meetings proving beneficial to keep us on track and discuss the project requirements. The R3 reading introduced us to organizing tasks within the team and implementing the ideas on Github's dashboards for tracking progress. R6 reading helped us use Git for our tasks with individual branches and then review each of our code to finally piece together each of our contributions to the master for submission.

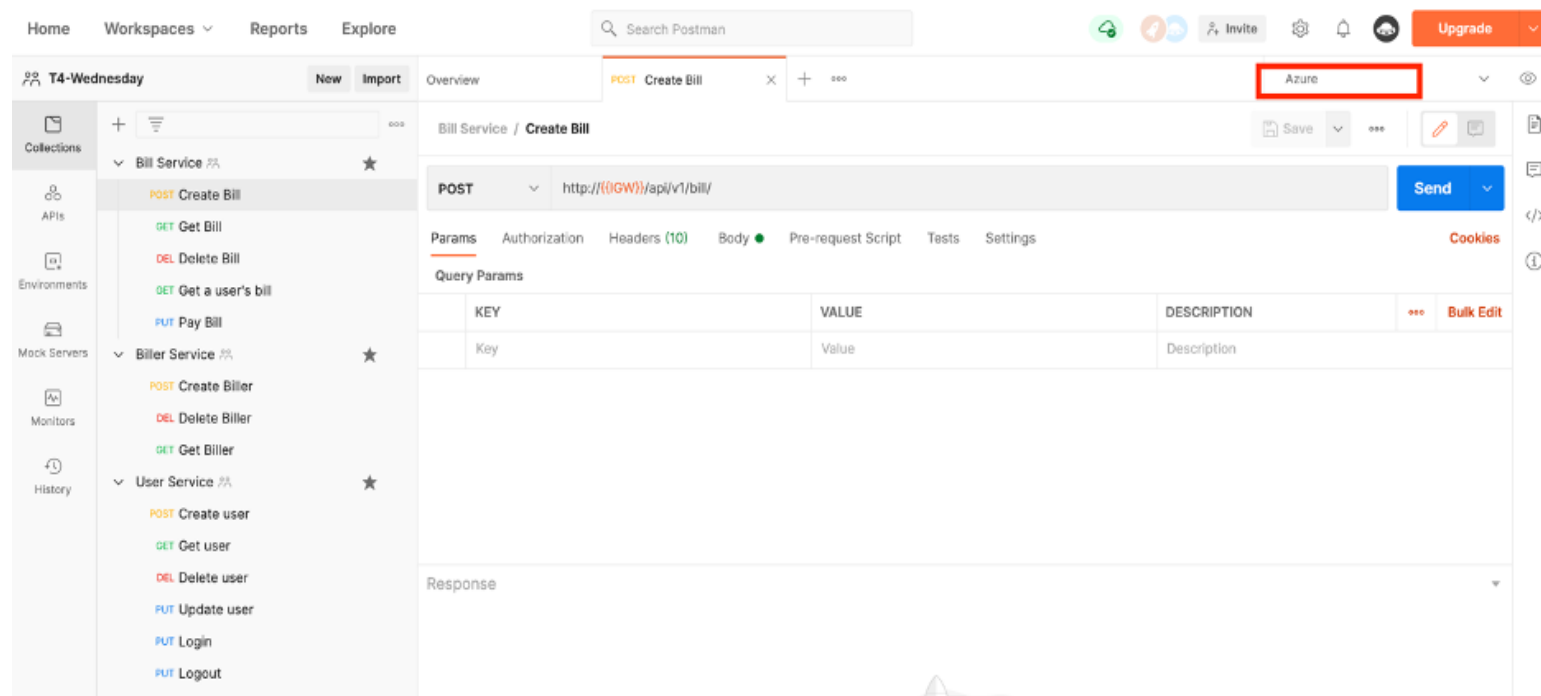
Our team performs well, adheres to the scrum principles, makes good progress, and handles tasks regularly. In our professional experience, we belonged to larger teams with a scrum master to supervise the progress. For this term project, the unavailability of

the product owner or scrum master added a bit more load on each of us, but we managed to take responsibility for our tasks and bring it together for the project. We believe that scrum success depends a lot on the team members if they actively participate in the meetings, collaborate well, etc. Consequently, we will ensure that each of us is proactive in his work and makes valuable contributions at every stage of the development lifecycle.

## Section 4 - Analysis

### Section 4.1 - Approach

1. We created and set up our team Postman workspace (T4-Wednesday) with all microservices (collections) and Request API's that are required for our application. We also created 2 different environments (Minikube, Azure) for testing and recording our API's through Gatling.



2. We set Postman proxy (localhost:8000) for sending requests to our API's. We recorded our basic application flow (Coverage) using Gatling recorder (HTTP Proxy mode) to create a scala file.

3. We took the generated script and ran the Coverage simulation. We analyzed the Coverage simulation results (report) for possible inferences and network analysis (refer to Section 4.2).
4. For Load Simulation, we edited the generated file with Azure cluster details and ramped up the users (50), requests/s and ran it for 30 minutes to see its behaviour. (refer to Section 4.2)

## Section 4.2 - Testing Analysis

### Planning:

- API methods in scope for testing:
  - Create Bill
  - Get Bill
  - Delete Bill
  - Get User Bill
  - Pay Bill
  - Create Biller
  - Delete Biller
  - Get Biller
  - Create User
  - Get User
  - Delete User
  - Update User
  - User Login
  - User Logout
- Testing scenarios:
  - Scenario 1 - Coverage simulation: **One active user** tests all APIs.
  - Scenario 2 - Load simulation: **50 active users, 20 API calls/sec, 30 minutes duration**
  - Scenario 3 - Load simulation with failure analysis:
    - Possible failover cases:
      - Istio Service Mesh
        - Circuit breakers
        - Retries
        - Fault injections
      - Cluster nodes failover



## Testing results:

### Scenario 1: Coverage simulation:

#### - Findings:

- It was possible to execute all the API operations in scope without errors. No HTTP failed requests occurred.
- “Create” operations: *Create Biller*, and *Create User* had a similar response time: 111ms and 106ms, respectively. However, in the case of “*Create Bill*”, the response time was 3064ms. We’ll wait for load simulation results to determine if this issue was a one-time situation or if it is necessary to review the implementation of the “*Create Bill*” method to identify potential optimizations.
- The remaining operations had a similar response time of 100ms on average, which is the expected response time defined by us for API responses (max 300ms):
  - “Retrieve” operations (*Get Bill*, *Get User Bill*, *Get Biller*, and *Get User*) response time: 100ms on average.
  - “Delete” operations (*Delete Bill*, *Delete Biller*, and *Delete User*) response time: 100ms on average.
  - “*Pay Bill*” operation: This PUT method had a response time of 104ms.
  - “*Update User*” operation: This PUT method had a response time of 143ms.
  - “Login/Logout” operations (*User Login* and *User Logout*): These PUT methods had a response time of 100ms on average.

### Scenario 2: Load simulation:

#### - Findings:

- “Create” operations (*Create New User*, *Create New Bill* and *Create New Biller*) had response times on average of about 2195 ms.
- “Retrieve” operations: The *Get* Requests had the best average response times of about 350 ms.
-

- “Pay Bill” operation: This PUT method had an average response time of 3497 ms, with zero failed requests.
- “Delete” operations (*Delete User*, *Delete Bill*, and *Delete Biller*): The three delete operations except the *Delete Biller* (212 ms) had high response times of about 2500 ms.
- As we can infer from the below graph that many KO (failed) requests were observed when we performed Load testing as compared to Coverage simulation. But, overall the OK requests are many and we can say that our cluster infrastructure (2 worker nodes) was not good enough to sustain that load. Given more nodes and capacity we could have minimized the failed requests.

more features with Gatling FrontLine

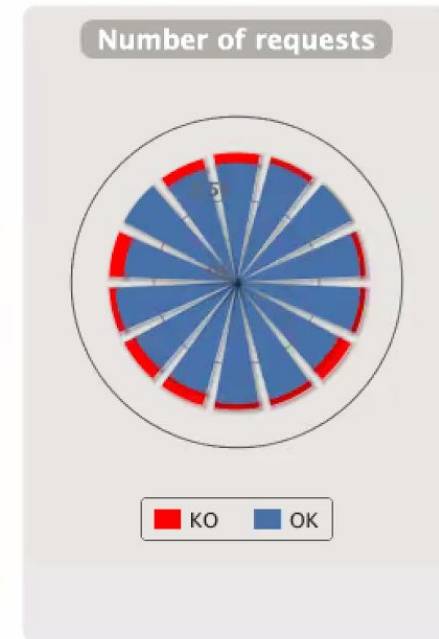
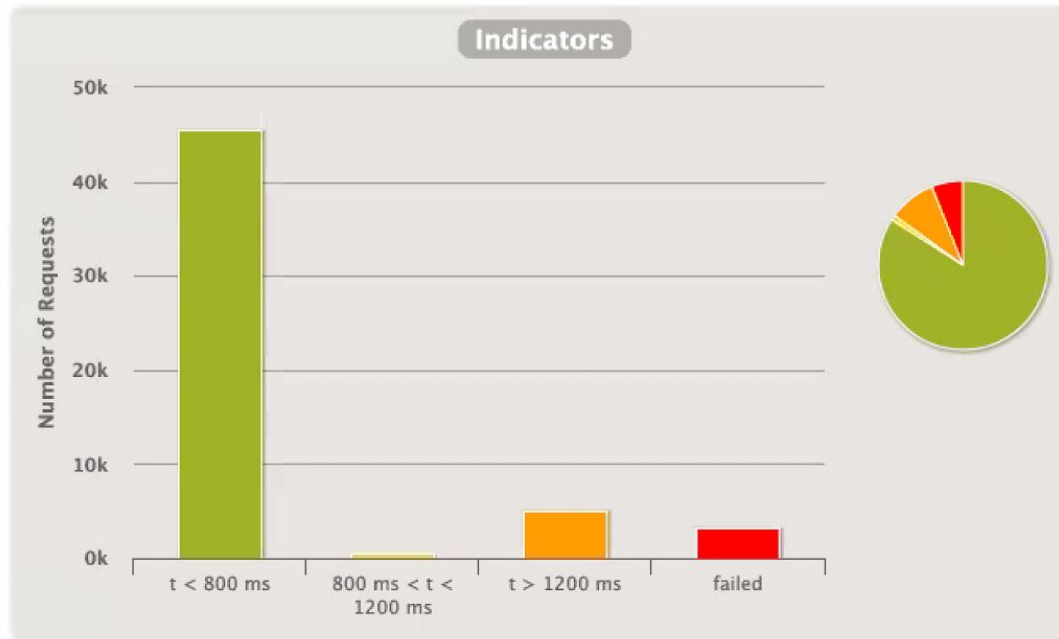


GLOBAL

▼ DETAILS

2020-12-04 23:22:22 -08:00, duration : 1800 seconds Load Simulation Test

### > Global Information



► STATISTICS

Expand all groups | Collapse all groups

STATISTICS <span>Expand all groups   Collapse all groups</span>													
Requests ^	Executions					Response Time (ms)							
	Total ↕	OK ↕	KO ↕	% KO ↕	Cnt/s ↕	Min ↕	50th pct ↕	75th pct ↕	95th pct ↕	99th pct ↕	Max ↕	Mean ↕	Std Dev ↕
Global Information	54385	51064	3321	6%	30.214	66	172	377	13195	26145	27487	1648	4984
Create User	3911	3600	311	8%	2.173	106	253	438	13452	26268	27338	2195	5927
Get User	3911	3735	176	5%	2.173	105	156	345	642	1191	3199	254	232
Log in User	3911	3911	0	0%	2.173	67	75	80	172	599	2513	98	124
Create Biller	3911	3737	174	4%	2.173	74	150	328	568	910	1891	235	180
Get Biller	3911	3738	173	4%	2.173	73	148	331	573	852	1808	230	172
Create Bill	3911	3493	418	11%	2.173	74	555	6712	26009	26385	27270	4873	7710
Get Bill	3879	3681	198	5%	2.155	74	156	333	527	1162	27293	348	1698
Get Bills for User	3863	3678	185	5%	2.146	74	154	330	410	856	26467	243	749
Pay Bill	3863	3502	361	9%	2.146	74	335	847	25997	26348	27262	3497	7125
Delete Bill	3863	3505	358	9%	2.146	74	237	439	25990	26378	27325	2529	6488
Delete Biller	3863	3666	197	5%	2.146	74	139	322	585	748	1351	221	157
Update User	3863	3414	449	12%	2.146	105	659	6733	26104	26433	27467	5501	8194
Logoff User	3863	3863	0	0%	2.146	66	75	81	203	530	1957	97	102
Delete User	3862	3541	321	8%	2.146	106	311	595	25988	26293	27487	2788	6526

- During the load simulation we found that our db-container was in an **Evicted state** during the run. As we can see that initially we had 17 pods running in total (all types) but after db-container pod was evicted, a new pod was initialized and now there are 18 total pods (See stats below line graph).

```

kubect1 get gw,deployments,pods
NAME                                     AGE
gateway.networking.istio.io/my-gateway 4d14h

NAME                                READY    UP-TO-DATE    AVAILABLE    AGE
deployment.apps/billcontainer        1/1      1              1            4d14h
deployment.apps/billercontainer      1/1      1              1            4d14h
deployment.apps/dbcontainer          1/1      1              1            4d14h
deployment.apps/usercontainer        1/1      1              1            4d14h

NAME                                READY    STATUS    RESTARTS    AGE
pod/billcontainer-5c869cf66f-d9jbx   2/2      Running   0           26m
pod/billercontainer-99c4447c8-2x4nf  2/2      Running   0           26m
pod/dbcontainer-6bf7b5ff77-4f7x9     2/2      Running   0           58s
pod/dbcontainer-6bf7b5ff77-swjn2     0/2      Evicted   0           26m
pod/usercontainer-5d44df8dcf-pkqbl   2/2      Running   0           26m
kubect1 -n termproject get svc
NAME                                TYPE        CLUSTER-IP    EXTERNAL-IP    PORT(S)    AGE
billcontainer                       ClusterIP    10.0.27.12    <none>         5002/TCP   4d14h
billercontainer                     ClusterIP    10.0.156.143  <none>         5001/TCP   4d14h
dbcontainer                         ClusterIP    10.0.149.147  <none>         5000/TCP   4d14h
usercontainer                       ClusterIP    10.0.79.13    <none>         5003/TCP   4d14h

```

Dashboard >

## Metrics

Azure Monitoring

+ New chart Refresh Share Feedback

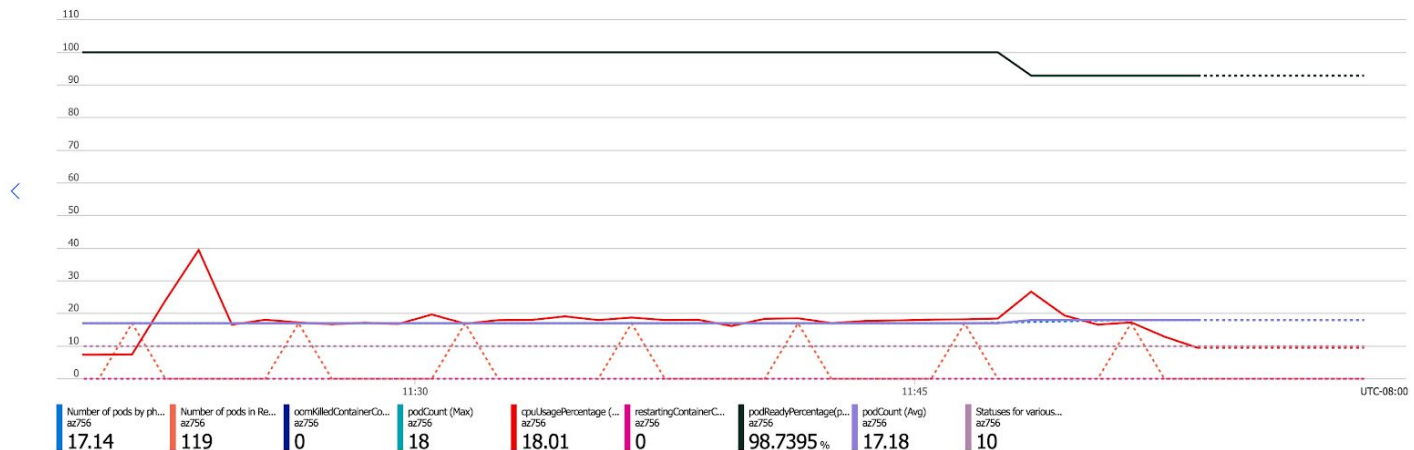
Local Time: 12/4 11:20 PM - 11:59 PM (Automa...

Add metric Add filter Apply splitting

Line chart Drill into Logs New alert rule Save to dashboard

You have unsaved changes to the chart. You can save the chart back to dashboard or pin it as a new chart to the dashboard.

az756, Number of pods by pha... Avg az756, Number of pods in Rea... Sum az756, oomKilledContainerCo... Avg az756, podCount, Max az756, cpuUsagePercentage, Avg az756, restartingContainerCo... Max az756, podReadyPercentage(... Avg az756, podCount, Avg az756, Statuses for various no... Avg



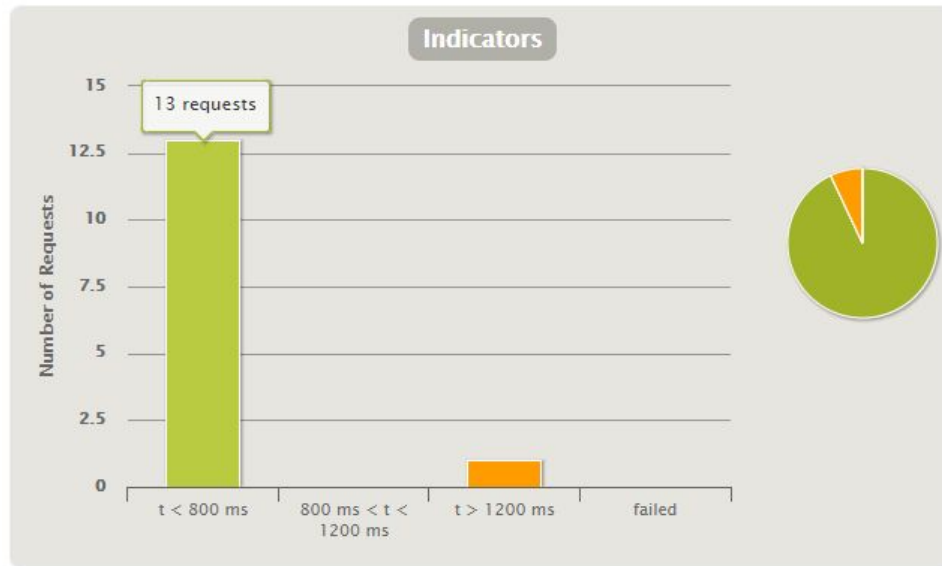
Azure metrics

### Scenario 3: Load Simulation with Failure analysis:

- **Deleting a pod (Azure Metrics):** We observed that after deleting a pod, k8s create a new set of pod to compensate for the down node and try to pick up the load again. This essentially works in a similar way how HDFS and other distributed systems work which separates the services and the underlying system.
- **Increased user load (150):** When we increased the user load for our application, we found out that few containers were in evicted state and k8s initialized new pods.
- **Re-route/ Shutdown isio service mesh** - Pending.

## Appendix A: Gatling statistics for the **coverage simulation**:

- Global information:
- Indicators:



- Statistics:



STATISTICS <span>Expand all groups   Collapse all groups</span>													
Requests ^	Executions					Response Time (ms)							
	Total ↕	OK ↕	KO ↕	% KO ↕	Cnt/s ↕	Min ↕	50th pct ↕	75th pct ↕	95th pct ↕	99th pct ↕	Max ↕	Mean ↕	Std Dev ↕
Global Information	14	14	0	0%	0.318	4	106	108	1165	2684	3064	312	764
Create Bill	1	1	0	0%	0.023	3064	3064	3064	3064	3064	3064	3064	0
Get Bill	1	1	0	0%	0.023	106	106	106	106	106	106	106	0
Delete Bill	1	1	0	0%	0.023	106	106	106	106	106	106	106	0
Get User Bill	1	1	0	0%	0.023	107	107	107	107	107	107	107	0
Pay Bill	1	1	0	0%	0.023	104	104	104	104	104	104	104	0
Create Biller	1	1	0	0%	0.023	111	111	111	111	111	111	111	0
Delete Biller	1	1	0	0%	0.023	100	100	100	100	100	100	100	0
Get Biller	1	1	0	0%	0.023	99	99	99	99	99	99	99	0
Create User	1	1	0	0%	0.023	106	106	106	106	106	106	106	0
Get User	1	1	0	0%	0.023	103	103	103	103	103	103	103	0
Delete User	1	1	0	0%	0.023	104	104	104	104	104	104	104	0
Update User	1	1	0	0%	0.023	143	143	143	143	143	143	143	0
User Login	1	1	0	0%	0.023	108	108	108	108	108	108	108	0
User Logout	1	1	0	0%	0.023	4	4	4	4	4	4	4	0